

# Oracle Database Performance Tuning Guide 阅读笔记

## Managing snapshots

eg:

```
begin
```

```
    dbms_workload_repository.create_snapshot();
```

```
end;
```

```
/
```

Note: a snapshot is created immediately on the local database instance with the flush level specified to the default flush level of TYPICAL. To view information about the newly created snapshot (and any existing snapshots), use the DBA\_HIST\_SNAPSHOT view.

eg:

```
begin
```

```
    dbms_workload_repository.drop_snapshot_range(low_snap_id=>22, high_snap_id=32, dbid=> 3310949047);
```

```
end;
```

```
/
```

Note: snapshots with snapshot IDs ranging from 22 to 32 are dropped immediately from the database instance with the database identifier of 3310949047. Any ASH data that were captured during this snapshot range are also purged.

eg:

```
begin
```

```
    dbms_workload_repository.modify_snapshot_settings(retention => 4, interval => 30, topnsql => 100, dbid => 3310949047);
```

```
end;
```

```
/
```

Note: the snapshot setting for the database with the database identifier 3310949047 are modified as follows:

The retention period is specified as 43200 minutes (30 days).

The interval between each snapshot is specified as 30 minutes.

The number of Top SQL to flush for each SQL criteria is specified as 100.

## Managing Baselines

eg:

```
begin
```

```
    dbms_workload_repository.create_baseline(start_snap_id => 270, end_snap_id => 280, baseline_name => 'peak baseline', dbid => 3310949047, expiration => 30);
```

```
end;
```

```
/
```

Note: a baseline is created on the database instance with the database identifier of 3310949047 with the following settings:

The start snapshot sequence number is 270.

The end snapshot sequence number is 280.

The name of baseline is peak baseline.

The expiration of the baseline is 30 days.

eg:

```
begin
```

```
    dbms_workload_repository.drop_baseline (baseline_name=> 'peak baseline', cascade => false, dbid => 3310949047);
```

```
end;
```

```
/
```

Note: the baseline peak baseline is dropped from the database instance with the database identifier of 3310949047 and the associated snapshots are preserved.

eg:

```
begin
```

```
    dbms_workload_repository.rename_baseline(old_baseline_name => 'peak baseline', new_baseline_name => 'peak Mondays', dbid => 3310949047);
```

```
end;
```

```
/
```

Note: the name of the baseline on the database instance with the database identifier of 3310949047 is renamed from peak baseline to peak Mondays.

eg: to display the summary statistics for metric values in a baseline period using the command-line interface, use the `SELECT_BASELINE_METRICS` function.

`dbms_workload_repository.select_baseline_metrics(baseline_name in varchar2, dbid in number default null, instance_num in number default null)` return `awr_baseline_metric_type_table` pipelined;

eg:

```
begin
    dbms_workload_repository.modify_baseline_window_size(window_size=>30, dbid=> 3310949047);
end;
/
```

Note: the default moving window is resized to 30 days on the database instance with the database identifier of 3310949047.

## Managing Baseline Templates

eg:

```
begin
    dbms_workload_repository.create_baseline_template(
        start_time => '2012-04-02 17:00:00 PST',
        end_time => '2012-04-02 20:00:00 PST',
        baseline_name => 'baseline_120402',
        template_name => 'template_120402',
        expiration => 30,
        dbid=> 3310949047
    );
end;
/
```

Note: a baseline template named `template_120402` is created that will generate a baseline named `baseline_120402` for the time period from 5:00 p.m. to 8:00 p.m. on April 2, 2012 on the database with a dbid of 3310949047. The baseline will expire after 30 days.

eg: to create a repeating baseline template using command-line

```
BEGIN
    DBMS_WORKLOAD_REPOSITORY.CREATE_BASELINE_TEMPLATE (
        day_of_week => 'monday', hour_in_day => 17,
        duration => 3, expiration => 30,
        start_time => '2012-04-02 17:00:00 PST',
        end_time => '2012-12-31 20:00:00 PST',
        baseline_name_prefix => 'baseline_2012_mondays_',
        template_name => 'template_2012_mondays',
        dbid => 3310949047);
END;
/
```

Note: In this example, a baseline template named `template_2012_mondays` is created that will generate a baseline on every Monday from 5:00 p.m. to 8:00 p.m. beginning on April 2, 2012 at 5:00 p.m. and ending on December 31, 2012 at 8:00 p.m. on the database with a database ID of 3310949047. Each of the baselines will be created with a baseline name with the prefix `baseline_2012_mondays_` and will expire after 30 days.

eg: To drop a baseline template using command-line

```
BEGIN
    DBMS_WORKLOAD_REPOSITORY.DROP_BASELINE_TEMPLATE (
        template_name => 'template_2012_mondays',
        dbid => 3310949047);
END;
/
```

Note: the baseline template named `template_2012_mondays` is dropped from the database instance with the database identifier of 3310949047.

#### Transporting Automatic Workload Repository Data

eg: to extract AWR data

```
@$ORACLE_HOME/rdbms/admin/awrextr.sql
```

and then enter the necessary parameter as prompted.

#### Loading AWR Data

eg: to load AWR data

```
@$ORACLE_HOME/rdbms/admin/awrload.sql
```

#### Using Automatic Workload Repository Views

DBA\_HIST\_ACTIVE\_SESS\_HISTORY

DBA\_HIST\_BASELINE

DBA\_HIST\_BASELINE\_DETAILS

DBA\_HIST\_BASELINE\_TEMPLATE

DBA\_HIST\_DATABASE\_INSTANCE

DBA\_HIST\_DB\_CACHE\_ADVICE

DBA\_HIST\_DISPATCHER

DBA\_HIST\_DYN\_REMASTER\_STATUS

DBA\_HIST\_IOSTAT\_DETAIL

DBA\_HIST\_SHARED\_SERVER\_SUMMARY

DBA\_HIST\_SNAPSHOT

DBA\_HIST\_SQL\_PLAN

DBA\_HIST\_WR\_CONTROL

#### Generating Automatic Workload Repository Reports

User Interface for Generating an AWR Report : Oracle Enterprise Manager

or

Generating an AWR Report Using the Command-Line Interface

The DBA role is required to run these scripts.

1. Generating an AWR Report for the Local Database

```
@$ORACLE_HOME/rdbms/admin/awrrpt.sql
```

2. Generate an AWR report on a specific database instance using the command-line interface

```
@$ORACLE_HOME/rdbms/admin/awrrpti.sql
```

3. Generating an Oracle RAC AWR Report for the Local Database

```
@$ORACLE_HOME/rdbms/admin/awrgprt.sql
```

4. Generating an Oracle RAC AWR Report for a Specific Database

```
@$ORACLE_HOME/rdbms/admin/awrgpti.sql
```

5. Generating an AWR Report for a SQL Statement on the Local Database

```
@$ORACLE_HOME/rdbms/admin/awrsqrpt.sql
```

6. Generating an AWR Report for a SQL Statement on a Specific Database

```
@$ORACLE_HOME/rdbms/admin/awrsqrpi.sql
```

7. Generating a Performance Hub Active Report Using a SQL Script

```
@$ORACLE_HOME/rdbms/admin/perfhubrpt.sql
```

#### ADDDM Analysis

An ADDM analysis can be performed on a pair of AWR snapshots and a set of instances from the same database.

#### Comparing Database Performance Over Time

User Interfaces for Generating AWR Compare Periods Reports

The primary interface for generating AWR Compare Periods reports is Oracle Enterprise Manager.

Generating an AWR Compare Periods Report Using the Command-Line Interface

Generating an AWR Compare Periods Report on the Local Database instance

```
@$ORACLE_HOME/rdbms/admin/awrddrpt.sql
```

Generating an AWR Compare Periods Report for a Specific Database

```
@$ORACLE_HOME/rdbms/admin/awrddrpi.sql
```

Generating an Oracle RAC AWR Compare Periods Report for the Local Database

```
@$ORACLE_HOME/rdbms/admin/awrgdrpt.sql
```

Generating an Oracle RAC AWR Compare Periods Report for a Specific Database

```
@$ORACLE_HOME/rdbms/admin/awrgdrpi.sql
```

Analyzing Sampled Data

Generating Active Session History Reports

User Interfaces for Generating ASH Reports

The primary interface for generating ASH reports is Oracle Enterprise Manager.

Generating an ASH Report Using the Command-Line Interface

to generate an ASH report on the local database instance

```
@$ORACLE_HOME/rdbms/admin/ashrpt.sql
```

to generate an ASH report on a specific database instance using the command-line interface

```
@$ORACLE_HOME/rdbms/admin/ashrpti.sql
```

Generating an ASH Report for Oracle RAC

to generate an ASH report for Oracle RAC

```
@$ORACLE_HOME/rdbms/admin/ashrpti.sql
```

## 10 Instance Tuning Using Performance Views

Instance Tuning Steps

1. Define the Problem
2. Examine the Host System and Examine the Oracle Database Statistics
3. Implement and Measure Change
4. Iterate 1, 2, 3.

To reduce I/O load caused by Oracle Database, examine the I/O statistics collected for all I/O calls made by the database using the following views

V\$IOSTAT\_CONSUMER\_GROUP : captures I/O statistics for consumer groups.

V\$IOSTAT\_FILE : captures I/O statistics of database files that are or have been accessed.

V\$IOSTAT\_FUNCTION : captures I/O statistics for database functions (such as the LGWR and DBWR).

eg:

```
select function_id, function_name from v$iostat_function order by function_id;
```

V\$IOSTAT : contains I/O statistics for both single and multi block read and write operations.

Identifying I/O Problems Using Operating System Monitoring Tools

eg:

```
sar -d
```

or

```
iostat
```

Identifying Network Issues

V\$IOSTAT\_NETWORK

Examine the Oracle Database Statistics

Setting the Level of Statistics Collection

STATISTICS\_LEVEL : the initialization parameter which controls all major statistics collection or advisories in the database.  
three levels:

BASIC

TYPICAL

ALL

Dynamic Performance Views Contain Wait Event Statistics

## V\$ACTIVE\_SESSION\_HISTORY

V\$SESS\_TIME\_MODEL : displays active database session activity, sampled once every second.

V\$SYS\_TIME\_MODEL : contain time model statistics, including DB time which is the total time spent in database calls.

V\$SESSION\_WAIT : displays information about the current or last wait for each session.

V\$SESSION : displays info. about each current session and contains the same wait statistics as those found in the V\$SESSION\_WAIT view.

V\$SESSION\_EVENT : provides summary of all the events the session has waited for since it started.

V\$SESSION\_WAIT\_CLASS : provides the number of waits and the time spent in each class of wait events for each session.

V\$SESSION\_WAIT\_HISTORY : displays information about the last ten wait events for each active session.

V\$SYSTEM\_EVENT : provides a summary of all the event waits on the instance since it started.

V\$EVENT\_HISTOGRAM : displays a histogram of the number of waits, the maximum wait, and total wait time on an event basis.

V\$FILE\_HISTOGRAM : displays a histogram of times waited during single block reads for each file.

V\$SYSTEM\_WAIT\_CLASS : provides the instance wide time totals for the number of waits and the time spent in each class of wait events.

V\$TEMP\_HISTOGRAM : displays a histogram of times waited during single block reads for each temporary file.

## System Statistics

V\$ACTIVE\_SESSION\_HISTORY : displays active database session activity, sampled once every second.

V\$SYSSTAT : contains overall statistics for many different parts of Oracle Database, including rollback, logical and physical I/O, and parse data.

V\$FILESTAT : contains detailed file I/O statistics for each file, including the number of I/Os for each file and the average read time.

V\$ROLLSTAT : contains detailed rollback and undo segment statistics for each segment.

V\$ENQUEUE\_STAT : contains detailed enqueue statistics for each enqueue, including the number of times an enqueue was requested and the number of times an enqueue was waited for, and the wait time.

V\$LATCH : contains detailed latch usage statistics for each latch, including the number of times each latch was requested and the number of times the latch was waited for.

## Segment-Level Statistics

You can query segment-level statistics through the following dynamic performance views:

V\$SEGSTAT\_NAME : lists the segment statistics being collected and the properties of each statistic.

V\$SEGSTAT : a highly efficient, real-time monitoring view that shows the statistic value, statistic name, and other basic information.

V\$SEGMENT\_STATISTICS : a user-friendly view of statistic values. In addition to all the columns of V\$SEGSTAT, it has information about such things as the segment owner and table space name. It makes the statistics easy to understand, but it is more costly.

## Implement and Measure Change

### Interpreting Oracle Database Statistics

#### Examine Load

Load-related statistics to examine include redo size, session logical reads, db block changes, physical reads, physical read total bytes, physical writes, physical write total bytes, parse count (total), parse count (hard), and user calls. This data is queried from V\$SYSSTAT.

#### Changing Load

#### High Rates of Activity

1. A hard parse rate of more than 100 a second indicates that there is a very high amount of hard parsing on the system.
2. Check whether the sum of the wait times for library cache and shared pool latch events is significant compared to statistic DB time found in V\$SYSSTAT. If so, examine the SQL ordered by Parse Calls section of the AWR report.

A high soft parse rate could be in the rate of 300 a second or more.

#### Using Wait Event Statistics to Drill Down to Bottlenecks

The most effective way to use wait event data is to order the events by the wait time. This is only possible if TIMED\_STATISTICS is set to true. Otherwise, the wait events can only be ranked by the number of times waited, which is often not the ordering that best represents the problem.

To get an indication of where time is spent, follow these steps:

1. Examine the data collection for V\$SYSTEM\_EVENT. The events of interest should be ranked by wait time.

- Alternatively, look at the Top 5 Timed Events section at the beginning of the Automatic Workload Repository report.
2. Look at the number of waits for these events, and the average wait time.
3. The top wait events identify the next places to investigate. It is usually a good idea to also have quick look at high-load SQL.
4. Examine the related data indicated by the wait events to see what other information this data provides. Determine whether this information is consistent with the wait event data.
5. To determine whether this theory is valid, cross-check data you have examined with other statistics available for consistency.

#### Table of Wait Events and Potential Causes

buffer busy waits : depends on buffer type, eg. waits for an index block may be caused by a primary key that is based on an ascending sequence. (V\$SESSION)

free buffer waits : slow DBWR (possibly due to I/O), or Cache too small. (Examine write time using operating system statistics. Check buffer cache statistics for evidence of too small cache.)

db file scattered read : Poorly tuned SQL, or Slow I/O system. (Investigate V\$SQLAREA to see whether there are SQL statements performing many disk reads. Cross-check I/O system and V\$FILESTAT for poor read time.)

db file sequential read : Poorly tuned SQL, or Slow I/O system. (same as the last one.)

enqueue waits (waits starting with enq:) : Depends on type of enqueue. (Look at V\$ENQUEUE\_STAT)

library cache latch waits : library cache, library cache pin, and library cache lock : SQL parsing or sharing. (Check V\$SQLAREA to see whether there are SQL statements with a relatively high number of parse calls or a high number of child cursors. Check parse statistics in V\$SYSSTAT and their corresponding rate for each second.)

log buffer space: Log buffer small, or Slow I/O system. (Check the statistic redo buffer allocation retries in V\$SYSSTAT. Check configuring log buffer section in configuring memory chapter. Check the disks that house the online redo logs for resource contention.)

log file sync : Slow disks that store the online logs, Un-batched commits. (Check the disks that house the online redo logs for resource contention. Check the number of transactions (commit + rollbacks) each second, from V\$SYSSTAT.

Additional Statistics (There are several statistics that can indicate performance problems that do not have corresponding wait events.)

#### Redo Log Space Requests Statistic

The V\$SYSSTAT statistic redo log space requests indicates how many times a server process had to wait for space in the online redo log, not for space in the redo log buffer. Use this statistic and the wait events as an indication that you must tune checkpoints, DBWR, or archiver activity, not LGWR. Increasing the size of the log buffer does not help.

#### Read Consistency

##### Table Fetch by Continued Row

You can detect migrated or chained rows by checking the number of table fetch continued row statistic in V\$SYSSTAT.

row migration

row chaining

\$ORACLE\_HOME/rdbms/admin/utlchnl.sql

#### Parse-Related Statistics

eg:

```
select name, value from v$sysstat where name in ('parse time cpu', 'parse time elapsed', 'parse count (hard)', 'CPU used by this session');
```

#### Wait Events Statistics

buffer busy waits

eg: to determine the possible causes, query V\$SESSION to identify the value of ROW\_WAIT\_OBJ# when the session waits for buffer busy waits.

```
select row_wait_obj# from v$session where event = 'buffer busy waits';
```

eg: to identify the object and object type contended for, query DBA\_OBJECTS using the value for ROW\_WAIT\_OBJ# that is returned from V\$SESSION.

```
select owner, object_name, subobject_name, object_type from dba_objects where dba_object_id = &row_wait_obj#;
```

If the contention is on the segment header,

eg: to find the current setting for free lists for that segment

```
select segment_name, freelists from dba_segments where segment_name = segmentname and segment_type = segmenttype;
```

Set free lists, or increase the number of free lists, or use free list groups.

If the contention is on tables or indexes (not the segment header)

Check for right-hand indexes. These are indexes that are inserted into at the same point by many processes

Consider using ASSM, global hash partitioned indexes, or increasing free lists to avoid multiple processes attempting to insert into the same block.

For contention on rollback segment header

If not using automatic undo management, then add more rollback segments.

For contention on rollback segment block.

If not using automatic undo management, then consider making rollback segment sizes larger.

db file scattered read

Inadequate I/O Distribution

Finding the SQL Statement executed by Sessions Waiting for I/O

eg: to determine which sessions are waiting for I/O

```
select sql_address, sql_hash_value from V$SESSION where event like 'db file%read';
```

Finding the Object Requiring I/O

eg: to identify the value of row\_wait\_obj# when the session waits for db file scattered read

eg:

```
select row_wait_obj# from V$SESSION where event = 'db file scattered read';
```

eg: to identify the object and object type contended for.

eg:

```
select owner, object_name, subobject_name, object_type from dba_objects where data_object_id = &row_wait_obj;
```

db file sequential read

direct path read and direct path read temp

